



<University Name>
<Department Name>

Summer Practice Report <Course Code>

Student Name:

Organization Name:

Address:

Start Date:

End Date:

Total Working Days:

Video Presentation Link (optional):

Contents

1	Introduction	3
2	Information about Project	4
2.1	Analysis Phase	4
2.1.1	Modules and Entities	5
2.1.2	Data	7
2.1.3	Design Constraints	8
2.1.4	Documentation	8
2.1.5	Conventions	9
2.2	Design Phase	9
2.3	Implementation Phase	11
2.3.1	Front-End Implementation	11
2.3.2	Back-End Implementation	14
2.3.3	DevOps Work	15
2.4	Testing Phase	15
3	Organization	16
3.1	Organization and Structure	16
3.2	Methodologies and Strategies Used in the Company	16
4	Conclusion	17
4.1	Appendices	17

1 Introduction

This document contains information regarding my internship in Siren Bilisim on summer of 2021. The document will consist of information about the organization, project I have worked on and my role in the project. Although its development process isn't perfectly aligned with software development lifecycle, the details regarding the project, namely PGMaster, will be organized as following:

- **Analysis:** In this section, I will be explaining the effort put into understanding and determining the requirements for the software.
- **Design:** In this section, the design choices that were made in accordance with the requirements will be demonstrated.
- **Implementation:** The development of the software is somewhat intertwined with the prior phases. However, explanations regarding the implementation specifically will be included in this section.
- **Testing:** This software's implementation involves thorough testing of its different modules. In terms of the timeline, the testing is not necessarily performed *after* implementation. Information regarding testing will still be written in this section of the document for sake of sticking with the outline demanded by the department.

The internship is performed in Siren Bilisim under the supervision of Koray Kocamaz. During that time, I have been mainly involved in web development activities. Before the internship, I have had developed various applications that utilizes web technologies. During the internship, I had chance to learn more about front-end development on top of back-end. The project mainly involves active participation of three developers. While I have performed my day-to-day development activities under supervision of Koray, the team member that I worked with the most was Derviş Mahmutoğlu from HAVELSAN. In that regard, I could partially claim that my internship project included the contribution of multiple organizations although HAVELSAN itself is not officially involved.

Although my involvement in PGMaster was not limited to the front-end, majority of the source code I have written would allow users interact with the underlying system. Due to that, I will be sharing extensive demonstration of the user interfaces.

On top of web development, the project revolves around databases, containers, and clusters. Domain knowledge about these entities have been also learned during this internship. I will be mentioning them in detail within relevant sections of the report.

2 Information about Project

In this chapter, I will be explaining technical details regarding the project. To provide better context, I shall be introducing the project first.

PGMaster is an administration tool that mainly intends to help database and system administrators. PGMaster provides an interface for the systems that utilize containers which are homes to databases. Using Docker containers, our system creates and manages clusters. These clusters offer capabilities like high availability and data integrity. With various configurations, databases can be set up with modes like master, standby and more. Intended users of the product is admins that manage systems with high amounts of transactions. Whether it is private or official organization, many admins perform high level database management operations using tools on command-line interface (CLI). PGMaster can be considered as an abstraction over such interfaces as it allows critical operations to be initiated, observed and completed using a graphical user interface (GUI).

PGMaster, although trivializing many tasks for admins, provides a CLI over the web application as well. This is achieved using WebSocket as well as various libraries which will be explained further in detail later on.

Without any due, I will get into the phases I have been through during my internship.

2.1 Analysis Phase

On my first day of the internship, after being given a computer and setting up my development environment, I have been introduced to Derviş. Prior to my involvement, he was the one and only contributor of the project. It was determined that I was going to implement front-end application of the system. After discussing the high level requirements of the system, it was clear to be that library that I was going to use for the project, namely Reactjs, was suitable. The project didn't have any document prior to implementation.

Before my involvement, Derviş had already implemented some of the back-end application modules which was using Java and Spring Boot framework. Being an experienced software developer (with 20+ years of experience), he was quite comfortable implementing Java and bash code. He also had extensive knowledge on PostgreSQL and database management systems in general. However, he needed a hand with the front-end application. There was only a JavaServer Pages (JSP) implementation with a form and some buttons written in HTML which were not sufficient for more sophisticated and dynamic nature of the system.

We have discussed the project briefly and I proceeded to learn Reactjs the following days. I will be telling about implementation related details that were discussed on

initial meetings on following sections. After understanding the general requirements of the system I have spent time analyzing the following:

- Candidate modules and entities.
- Nature of the data to be processed.
- Design related constraints.
- Documentation needs.
- Conventions to be followed during implementation.

I will try to explain considerations I have had for each of these items.

2.1.1 Modules and Entities

Before jumping into these items, it seems necessary to talk about PGMaster modules and entities. The modules can be listed as follows:

- **PGManager:** Even if the name suggests something similar to PGAdmin, a web-based GUI tool that is commonly used to interact with PostgreSQL databases, it is actually the back-end of PGMaster application. It is written in Java (using Spring Boot framework) and it has two main purpose: Performing CRUD operations on pgm database, the database that contains “meta” data regarding the application itself and manage operations that affect PGMaster entities. These entities will be explained below.
- **Grafana:** Grafana is a multi-platform open source analytics and interactive visualization web application. It provides charts, graphs, and alerts for the web when connected to data sources. Our web application’s front-end is intended to be embedded into one of Grafana’s custom panels. Within the context of PGMaster, this is its most relevant purpose. Other than that, various services and endpoints that works on the container of PGMaster provide the data Grafana shows. As understood, the application PGMaster is also containerized. The content of this container will also be mentioned later on the document.
- **WeTTY:** WeTTY (Web + TTY) is an npm module that makes it possible to have terminal access in browser over HTTP(S). As mentioned before, PGManager is intended to offer necessary tools that would help advanced users to perform various actions on CLI. We have decided that using a library for this purpose would be beneficial over implementing our own module. It basically utilizes WebSocket protocol to establish an SSH connection. In PGManager, we intend to offer access to PG hosts and containers using this library.
- **PG-Web:** This is the application I have implemented from the beginning using Typescript and Reactjs. It basically provides a GUI to PGManager users where

they can perform various actions such as creating and initializing PGContainers, taking backups and restoring them, monitoring clusters and their nodes' status etc. PGMaster has over 100 such operations. While most of them needs only one parameter to be provided, some of them require more sophisticated data to be provided. For sake of simplicity, only some portions of these operations will be explained on the following sections of the document.

These four modules are the main “pillar stones” of a PGMaster ecosystem. Organizations that have systems with multiple databases which needs to deal with high volumes of transactions while maintaining high availability and also security and integrity of their data may unfortunately not have admins that are capable of performing operations that were mentioned above trivially. These modules intend to offer tools that are necessary to maintain, operate and administrate these systems whether they have advanced database administration skills or not.

I have mentioned that PGMaster have different entities. I will be listing what they are and what purpose do they serve below:

- **PGContainer:** As mentioned briefly before, databases within PGMaster ecosystem lives inside Docker containers. These databases have different configurations depending on their cluster modes. These modes can be simple listed as (I am not creating another list!) master, standby and replica. Intuitively, these names indicate whether a container contains the “main” database or its “copies”. Depending on the status of a master database (and its container), various operations can be triggered using PG-Web. These containers come together and create a cluster, more specifically PGCluster. It is worth mentioning that PGMaster itself is a container as well. This container is intended to contain modules mentioned above. While the main benefit of using a container is to be able to deploy PGMaster on an organization easily, other benefits of using a container could also be considered as the added advantages.
- **PGHost:** Naturally, a host is required to run Docker engine and PGHost is the host that houses these containers. While PGHost mainly indicate a dedicated, physical server, it can also be a virtual private server (VPS). These hosts houses PGContainer and their availability directly impacts health of a PGCluster. Naturally, a PGContainer may have nodes that lives inside different PGHosts. I think for the scope of this document, this should suffice explaining what a PGHost is.
- **PGCluster:** PGCluster, to be brief, is the combination of multiple PGContainers. Its nodes (PGContainers) may or may not live inside different PGHosts. Containers alone do not provide hi(gh availability and clusters are utilized to provide it. PGClusters are shown with a topology inside PG-Web. Entity management view, which will be explained later on, allows users to configure their nodes inside a PGCluster.
- **PGBackupServer:** As understood from the name, it is simply a server that contains backup files of a PGContainer database. Backup data is transferred to

this entity on desired intervals and when the need rises, this entity provides the data that needs to be recovered.

These four entities were the main entities that were necessary to have a functional PGMaster ecosystem. I have spent first two weeks creating views to create, update and delete show these entities. The main benefit of using Reactjs over regular HTML/CSS/JS stack for the GUI was to be able to render these entities optimally and dynamically without having to update the page. While Javascript makes a regular web page dynamic to some extent, Reactjs allows users to manipulate data and have Document Object Model (DOM) to be synchronized seamlessly. As stated on the official documents of Reactjs, Web Components and Reactjs are complementary, i.e., they solve different problems. While Web Components provide encapsulation of reusable components, Reactjs provides declarative implementation. So it is safe to say that real benefit of using Reactjs is felt by developers, rather than users. I don't mean to go out of scope of this document and explain benefits of React, so I will keep it brief. In PGMaster context, React allows us to utilize having declarative implementation, meaning the developer can set the "rules" for a view and the view updates automatically with the rules provided by the developer without needing to refresh the page user interacts. I have already mentioned that PG-Web, the module I have implemented is embedded into Grafana. It goes without saying that user should be able to observe the changes on PGMaster GUI without needing to refresh the web page, especially considering that Grafana already allows this for other views than PGMaster. In that sense, usage of a library such as Reactjs was vital during the implementation.

2.1.2 Data

The data that circulates on PGMaster could be divided into two:

- **Entity Data:** PGMaster has its own database to contain details regarding entities. While adding/updating/deleting these entities, a PostgreSQL database (pgm) is used. PGManager (the back-end) has numerous endpoints (listHosts, listContainers etc.) that feed the front-end application PG-Web. Entity data is fetched to render initial view of PG-Web and it is updated performing various CRUD operations on the said database. When a user updates entities, the change is immediately reflected on PG-Web. This ensures users to have the latest updates on PGMaster to be synchronized on the GUI they are using. I have implemented the GUI in a way that a user can have sufficient feedback and information whether operations they have initiated were successful or not.
- **Operation Results:** While most of the configuration a user can make is intended to be performed once or twice, the operations on a PGMaster entity is expected to be more repetitive. Depending on the status of a PGMaster entity, these operations may fail or succeed. I have implemented a view where users may see the results of the operations they have initiated. We have realized that HTTP status codes wouldn't suffice since they don't provide information regarding what went during

the execution of a PGMaster operation. Then Derviş decided that we can use slightly more sophisticated response objects, namely Rx, to provide user more than just a status code. These objects contain fields like error code, command, result text, hint, execution time, container name, host name, cluster name, path, details and more. I have implemented various views that allow users to view these. I'd like to mention a detail regarding the implementation of these views. When a user triggers a command to execute, this command have potential to trigger other commands. These commands may or may not fail depending on a lot of factors (such as permissions, status of the target entity etc.) and user should be able to identify when did an operation has failed with precision. I have offered two different alternatives for users. Firstly, users may see these commands and their results with indentation on a table. As an alternative to this, users may also see a command and its "child" commands in a tree-like view. Upon discussions with Derviş, we have decided to use the first, as it would be practically impossible to show the results of these commands' results on a tree view, considering Grafana can only provide so much space.

2.1.3 Design Constraints

As mentioned before, PG-Web is intended to be shown among other views Grafana has. I must admit that I have faced the challenge of "unknown requirements". To elaborate, PGMaster is intended to be used by "power users" that have access to hardware such as 4K monitors. However, Derviş occasionally has tested the views on his personal mobile device and this forced me to consider another constraint which was being responsive. Fortunately, the library I have used (Semantic UI React) provided predefined solutions which allowed users to interact PG-Web seamlessly. Prior to my internship at Siren Bilisim, I have considered web development rather trivial. However, I have learned that implementing GUIs that just work on different devices is indeed a challenge where you have to consider countless edge cases. I am not sure how viable the end product will be on different mobile devices (tablets and mobile phones running Android or IOS) but I have spent decent amount of time on making our Uproduct work on different devices. Unfortunately, Siren Bilisim didn't have any UI/UX designer which was assigned to work on this project, so I had to come up with my own designs and implementations.

2.1.4 Documentation

Having had an introduction to software requirements specification (SRS) and software development lifecycle in general in our Software Engineering course, I usually try and understand the needs of any development activity prior to actual work. Aiming to follow the same trend here, I must admit that I was mildly disappointed, since although we didn't have a strict timeline to deliver the proof of concept (PoC) version of the application, the project was solely maintained and developed by one person. This should have -expectedly- obscured the need for any form of documentation.

I have actively communicated with my supervisor (Koray) and colleague (Derviş) on

the subject and I was given permission (and responsibility) to write SRS for the project. I did not start writing the document right away, because it took me a while to fully grasp the overall needs and constraints of the system and the rest of the requirements naturally needs further discussion with stakeholders and the potential users. I still had chance to write some of the document by myself and it will be attached to the report. Final version of the SRS is expected to be completed after it is presented to The Ministry of Health.

2.1.5 Conventions

Reactjs, the library that I have used extensively during the implementation of PG-Web is a library that was open-sourced in 2013. Since that time, with the help of community (that consists of developers from all around the world) Facebook switched from class-based components (and its implementation) to function-based components. For sake of not getting into too much details, I will skip how it happened, but I adopted the function-based component implementation approach. With the additions of “hooks”, Reactjs eliminated the need for most of the boilerplate code and I made use of this approach.

2.2 Design Phase

As explained on prior sections of the document, user interface-wise, the design was mostly determined on-the-go, that is, I have created the views using Semantic UI React and depending on the feedback I have received from my colleague Derviş.

When it comes to system design, I have adopted the design that was initially proposed and partially implemented by Derviş. Our application, although mainly relies on operating system level operations which were invoked by back-end application PGManager, is an MVC application in its essence. Using RESTful API to manage network requests, I have took part in development of back-end code in order to comply with the standards better.

It goes without saying that the actual layout the source code of the front-end and back-end applications are also relevant. This project, although having financial concerns and economic value, is also intended to be an open-source project in the long run. Having that in mind, I have tried to refactor the code base once I have reached certain milestones. For example, the source code for various React components would be on the same file, sometimes with poor readability. Since the company had limited man power when it comes to Reactjs, I have joined communities like Reactiflux and communicated with different developers around the world. This community actively communicates on their Discord server. The server has different channels where people ask and answer questions ranging from high-level to specific errors and issues with a React application. On my third week of the internship, I have posted on “code-review” channel and I was fortunate enough to receive feedback from a more experienced developer. Upon his feedback, I refactored about 40% of the codebase. To provide my development activities briefly without giving too much details, I will be sharing the commit history on appendices.

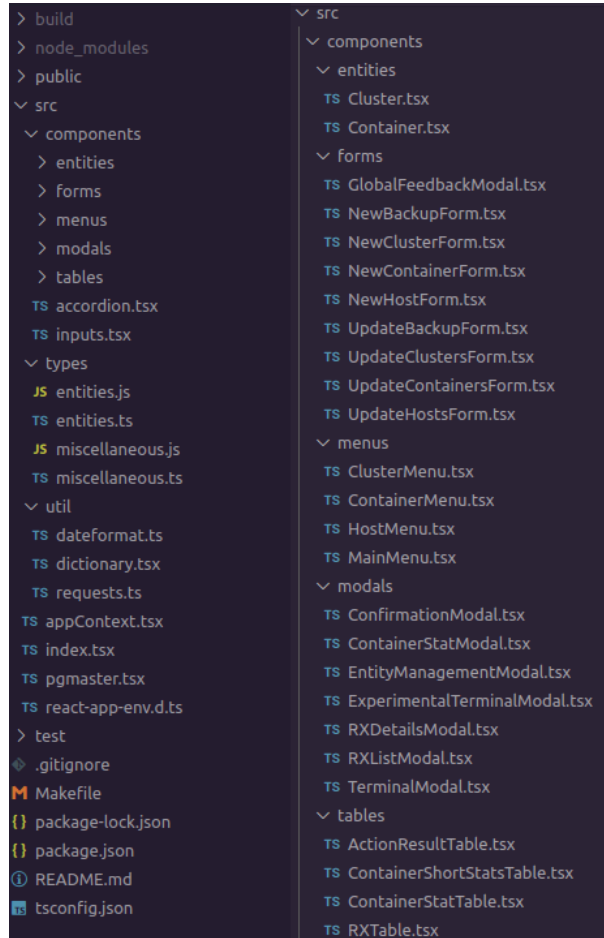


Figure 2.1: File structure of PG-Web

Considering the project hierarchy is more relevant to the *design* the latest version of the project (PG-Web, to be specific) can be seen above on Figure 2.1:

As mentioned before, PGMaster has its own database pgm to save entity related data. This database consists of four tables, pg.container, pg.host, pg.cluster and pg.backup. These table have their records updated with two main constaint. Regular SQL-related constraints and their dependency on the success of the operations that relates to them. In other words, when a user triggers a PGMaster operation from the GUI, it causes PGManager to execute various projects. While some of them utilize docker engine to start/stop/restart containers, some of them perform more sophisticated operations such as performing a HotStream Replica (with Delta Store), Wal Replica (with Delta Store), Diff backup (delta last full backup) etc. This database has a pretty straightforward schema which can be view on Figure 2.2 below:

pg_container	pg_host	pg_cluster
<ul style="list-style-type: none"> name varchar(50) cluster_name varchar(20) description varchar(50) enable_archive_mode boolean enable_audit boolean enable_auto_fail_over boolean enable_auto_recover boolean enable_cluster boolean enable_cron boolean enable_exporter boolean enable_pgagent boolean host_name varchar(50) image_name varchar(50) locale varchar(20) mode varchar(10) path varchar(50) port integer shell_access boolean wal_size integer fail_order_sequence integer 	<ul style="list-style-type: none"> name varchar(20) default_gateway varchar(20) description varchar(30) ip varchar(15) password varchar(30) port integer shell_access boolean ssh_public_key varchar(255) username varchar(30) 	<ul style="list-style-type: none"> name varchar(20) backup_name varchar(50) description varchar(50) enable_auto_fail_over boolean enable_auto_rebuild boolean enable_proxy boolean read_write_port integer readonly_port integer

pg_backup
<ul style="list-style-type: none"> name varchar(50) description varchar(255) ip varchar(15) key varchar(255) path varchar(255) port integer secret varchar(255) type varchar(255)

Figure 2.2: DB Schema of PGManager

2.3 Implementation Phase

2.3.1 Front-End Implementation

I had mentioned that I had prior front-end web development using HTML, CSS and Javascript. However, I didn't know Reactjs and I've spent approximately one week getting familiar with the library. Reactjs is a Javascript library that was priorly used by Facebook internally and later open-sourced around 2013. Since then the library have adopted different standards with the help of the community. I usually prefer sticking to official documents when I am learning a new technology. I have spent trying to understand what does Reactjs offers and what were the reasonings behind the changes that were added to it over the years. The main change that was implemented to the library was the addition of hooks. To be brief, initially Reactjs used class based components where you implemented a class with props and states, then adding bunch of bindings and other boilerplate code. From a developer's perspective of view, this was cumbersome. Recent versions of Reactjs encourage usage of hooks. Some of them are useState, useEffect, useContext, useMemo, useCallback and more. The library allows addition of custom hooks which modify the state of components that are widely used in Reactjs applications.

The main challenge I have faced learning and implementing Reactjs was that my prior experience with front-end web development involved writing source code that were mostly imperative and some functional paradigm. Reactjs, however, in addition to these, introduces wide usage of declarative paradigm where developer provides the rules for a

component and the component renders desired views respecting the rules provided by the developer. Most of the optimization is provided by the library out-of-the-box which saves a lot of unnecessary renders. With effective usage of proper hooks and determining props and states carefully also yields a responsive, fast view. As mentioned above, PGMaster is no less active when it comes to flow of data and it is easily comparable to modern social media web applications such as Facebook. Just like a user may modify the view of, say, chat box. Similarly, PGMaster is expected to render PGContainers and other relevant entities with their own Reactjs components considering their current status. The data that determines which entities and their respective React components to be rendered is synchronized with the data that is fed by PGManager, allowing users to see the results of their actions and their effects which is further determined by the operating system and Docker Engine.

I have made the mistake of refreshing the view using setInterval method of Javascript on my initial iterations which was conceptually against the way Reactjs is supposed to work. PG-Web is expected to trigger re-render without having to fetch the data that were updated by a user form, for example. To overcome this problem, I have learned about application-wide state solutions. One of the most common solutions the community offers was Redux, but I didn't want to use it for sake of simplicity and avoid going through the steep learning curve of the library. It turns out the main contributors of Reactjs were also aware of the issue at hand when it comes to store complicated data and they came up with "context". They added it with its own hook (in React 16) which enabled developers to set initial data, update it using various sources (user input or from a new work source) and make this data available to React components. The main benefit of using this feature comes from the elimination of prop drilling. Prop drilling is the method that were widely used prior to context would require developer to pass data from all the way down to the child components from the root component and that would be problematic especially when the source code needed refactoring. Reactjs provides a strong hierarchical component structure with the help of its composition capabilities.

As mentioned on prior sections, I have had the following data:

- Clusters
- Hosts
- Containers
- BackupServers

Although this data doesn't look too complex, I have noticed that every change in these entities' respective data structures, it would trigger a re-render on all the application. After reading Kent C. Dodds fantastic articles about [Application State Management with React](#) and [How to use React Context effectively](#). I came to realize that I was being too eager to use the new feature without making sure it was the best solution. When it comes to simple data such as "darkMode" which requires nothing but a boolean variable, React's Context really shines. As known, in today's web and mobile applications the option to have a dark/light mode is quite common. In such use cases, the application

needs to spread this data to a lot of nodes in its component tree. In such cases, a toggle that changes this variable needs to be reflected on all the components and trigger a re-render.

Long story short, I have realized that I can just use prop drilling to spread my entity data among components that need it. With prop drilling, there is a known trade-off between development time and application performance. In other words, although prop drilling means passing a variable from parent components to child components multiple times, it works seamlessly with only necessary amount of re-renders. I ended using this for entity data while using the new and shiny `useContext` hook for the `darkMode` which I already implemented on PG-Web.

Before proceeding with the rest of my implementation related activities, I'd like to comment a little about Typescript. Typescript is an open-source library that adds strongly typed language features into Javascript. This allows developers to avoid common pitfalls of using dynamically typed languages. Many of the errors that could be faced on run-time can be detected while “transpiling” with Typescript. Typescript basically transforms the code the developer created into a Javascript code in the end. However, it makes the code much safer with addition of types and interfaces. Without explaining too much about the language, I will share some of the types or interfaces I declared implementing PG-Web. These declarations allowed me to further utilize the capabilities of the text editor that I used by the help of language server, namely IntelliSense.

As seen on Figure 2.3 Typescript allows composition of basic primitive types. On top of the figure, interfaces (not to be confused with the other common usages of the word in computer science) for entities are declared. The middle interface makes sure that proper render props are passed to the OneMasterTwoSlavesTopology component. And on the bottom, the common usage of React's useContext hook can be viewed. I have used about 40 different object or method interfaces throughout the development PG-Web which further improved the speed of debugging and testing itself.

2.3.2 Back-End Implementation

Although my development activities mainly consisted of front-end development, following third week I felt comfortable enough with the business logic and participated on the development of the Java Spring Boot application. While sometimes work of Derviş was blocked by my implementation, there were days my work was blocked by his. I hopped in and modified some of the controller methods when needed. Furthermore, we have purchased three VPS from Contabo (German cloud services provider) and deployed our applications there. I have took active role during that process, but I'd like to tell about them on a separate section below. The reason I mentioned it here is that, we have encountered -supposedly- CORS errors while deploying. Prior to cloud deployment, Java and React application were running on localhost, eliminating most of the security

```
export interface IContainer {
  name: string;
  hostName: string;
  port: string;
  path: string;
  clusterName: string;
  description: string;
  imageName?: string;
  enableArchiveMode?: boolean;
  enablePgagent?: boolean;
  enableAudit?: boolean;
  enableCron?: boolean;
  enableExporter?: boolean;
  enableCluster?: boolean;
  enableAutoFailOver?: boolean;
  enableAutoRecover?: boolean;
  locale?: string;
  defaultGateway?: string;
  walSize?: string;
  shellAccess?: boolean;
  mode?: string;
}

export interface IHost {
  name: string;
  ip: string;
  port: number;
  description: string;
  username: string;
  password: string;
  sshPublicKey?: string;
  defaultGateway?: string;
  shellAccess?: boolean;
}

export interface ICluster {
  name: string;
  masterContainerName?: string;
  description: string;
  backupServer?: string;
  enableAutoFailOver?: boolean;
  enableAutoRebuild?: boolean;
  enableProxy?: boolean;
  readOnlyPort?: number;
  readWritePort?: number;
}

interface IOneMasterTwoSlavesTopologyProps {
  dispatch: React.Dispatch<PGMASTERACTION>;
  mastercontainer: IContainer;
  masterContainerHost: IHost;
  slaveOneContainer: IContainer;
  slaveOneContainerHost: IHost;
  slaveTwoContainer: IContainer;
  slaveTwoContainerHost: IHost;
}

const OneMasterTwoSlavesTopology = (props: IOneMasterTwoSlavesTopologyProps) => {
  const { darkMode } = useContext(ThemeContext);
  const updateXarrow = useXarrow();
}
```

Figure 2.3: Various interfaces that force type checking.

related concerns. However, after deploying on publicly available server, we have had issues where relevant configuration were needed. For that, I have refactored source code on Java where security configuration was being implemented. It was basically setting allowed origins and allowed REST methods (such as GET, POST, DELETE). While editing that code, I also edited some of the controller methods to better comply with RESTful API conventions. For example, the source code for deleting entities were using GET methods, I changed them into DELETE methods and changed the status code it would return on different cases. I have also had chance to refresh my memory on the term “idempotency”.

Other than changing controller and security configuration source code, I have updated some of the testing codes. Since it is much harder to initialize databases with the entity data we need for testing, we utilized Java for that. We had data initializer tests. These also acted as migration tools when we needed to update our entity schemas. When Derviş pushed a commit to back-end application and asked me to pull changes, I would simply drop my tables and run data initializers. This would add the necessary data to development database we used. Then I was able to run my react application and I would have the updated fields provided by back-end API. It goes without saying that I used various tools (RESTED the Firefox addon, Postman, even curl) to test PGManager’s endpoints. We have tried to enable Swagger to trivialize API documentation and testing processes, however it took too much time to set up and due to limited time we had before demo, we skipped it.

2.3.3 DevOps Work

Working with Docker containers and developing on Linux machines, it was inevitable to not get into DevOps related work in Siren. As I mentioned before to have a more realistic environment during development, we deployed our applications on cloud. I was able to run both PG-Web and PGManager locally, most of the operations PGMaster offers runs Docker engine and executes various commands that require extensive configuration. For that reason, I had to develop front-end application with “dummy” responses from the controller on back-end. Once we deployed on the server, I was able to receive more realistic responses, namely RX objects. These objects contained more details regarding the fate of an operation.

While deploying our applications, another point where I hit the wall was to configure Nginx. We used Nginx to utilize reverse proxy. As I have listed above, we have 4 different modules, all running a different processes on different ports. I chose the path of static hosting for the web application which left me with three applications to set up on Nginx.

2.4 Testing Phase

3 Organization

3.1 Organization and Structure

3.2 Methodologies and Strategies Used in the Company

4 Conclusion

4.1 Appendices